

Adaptive Write-Update and Write-Invalidate Cache Coherence Protocols for Producer-Consumer Sharing

By Bangjie Liu, Hao Li

Producer-Consumer Sharing

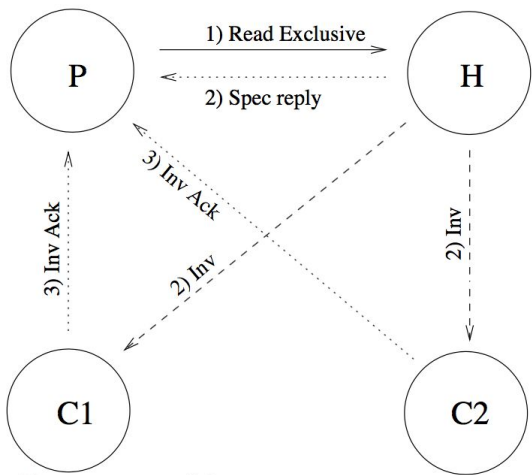
How to make it efficient from architectural perspective ?

Problem

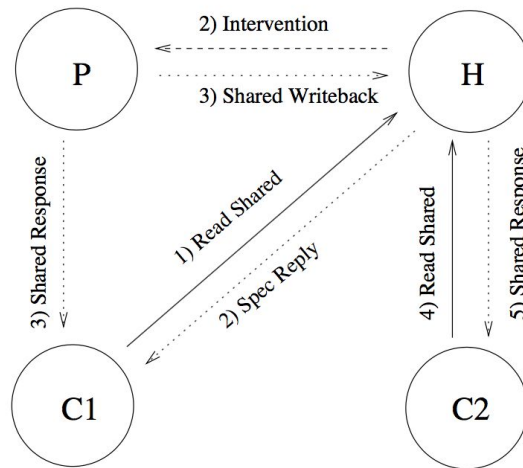
Typical cc-NUMA systems use directory-based write-invalidate protocols^[1], which is inefficient for producer-consumer sharing.

- Remote misses: modern processors don't support capability to push data to processor caches
- Communication traffics: 3 hops, write contention lead to ping-ponging

Producer side:



Consumer side:



[1] Cheng, Liqun, John B. Carter, and Donglai Dai. "An adaptive cache coherence protocol optimized for producer-consumer sharing." *High Performance Computer Architecture, 2007. HPCA 2007. IEEE 13th International Symposium on.* IEEE, 2007. * This slide uses example figure 1 from [1]

Inspiration

Bus-based write-update protocol is efficient for producer-consumer sharing.

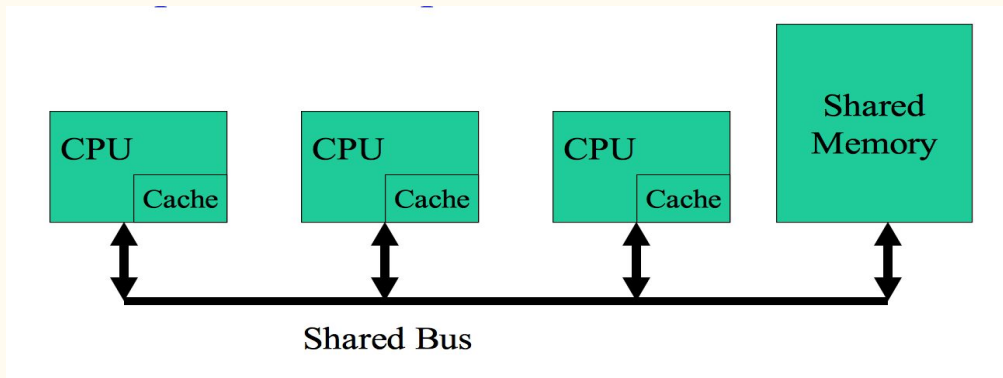
- Bus transactions are visible to all processors
- When local cache is updated by producer, the new data is broadcast to all consumer cache copies.

Benefits (for P-C sharing):

- Reduce communication traffics
- Reduce remote misses

Drawbacks (for general):

- Pattern dependent
- Bus bottleneck



[1] Prof. Nathan Beckmann, lecture slides "740: Computer Architecture Memory Consistency&Cache Coherence", Carnegie Mellon University

[2] Sundararaman Nakshatra, Architecture (EECC551) slides. "Cache Coherence Protocol", <http://meseec.ce.rit.edu/551-projects/fall2010/1-3.pdf>

* This slide uses the example figure from adapted lecture slide by Ian Watson, "Cache coherence in shared-memory architectures", University of Manchester, <http://www.cs.utexas.edu/~pingali/CS377P/2017sp/lectures/mesi.pdf>

Can we do better?


Adaptive Cache Protocol!

Solution


Directory-based adaptive cache coherence protocol

- Hardware support for cache → cache communication.
- Producer-Consumer sharing pattern detector.
- Capability to switch from write-invalidate to write-update.

Producer Side (on write)



Consumer Side (on read after copy invalidated)



Benefits

For P-C sharing:

- Reduce communication traffics: 3 hops \rightarrow 2 hops
- Reduce remote misses: cache \rightarrow cache

For general:

- Less bandwidth
- Only stakeholders retain copies

Potential Issues

- **Write propagation**
 - Hardware support, software test.
- **Write serialization**
 - Verify sequential consistency. (Murphi)

Goals

Goals

75%

- Implement a cache simulator and testing tools.
- Evaluate performance of directory-based cache-invalidate protocol on typical producer-consumer apps.

100%

- Implement proposed adaptive cache coherence protocols and producer-consumer sharing pattern detector.
- Evaluate the performance on producer-consumer apps.

125%

- Synchronize app threads to observe more accurate cache access events (eliminate pintool side effect).

Cache Simulator

Cache Simulator


We now have a highly customizable cache simulator. (based on open source Dynamic Cache Simulator for Shared Address Space (SAS))

- MSI (modified, shared, invalid)
- Directory-based
- Write-back-allocate

Representative Application

Representative Application

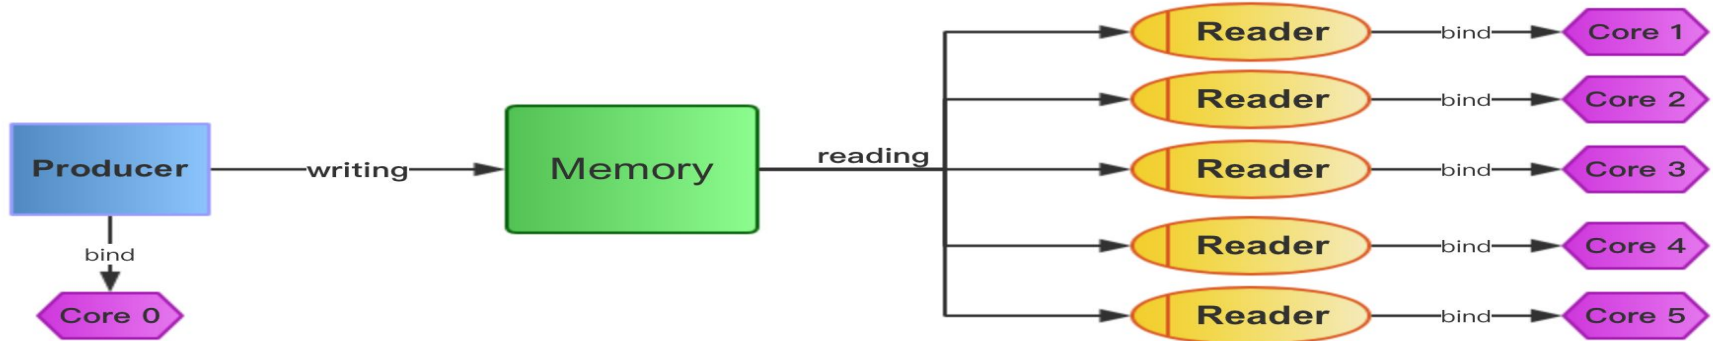
1. One producer keeps writing to a shared memory location while there are several readers keep reading from that particular location.
2. $(\# \text{physical cores} - 1)$ consumers



Representative Application

Potential issue: the OS can schedule threads among different cores, which will disturb cache.


Solution: bind threads explicitly to their designated cores



Representative Application

Potential issue: binding could happen after threads start running because binding takes time.

Solution: each thread sleeps for a certain amount of time before executing.



Result Analysis

Result Analysis

Experiment environment:

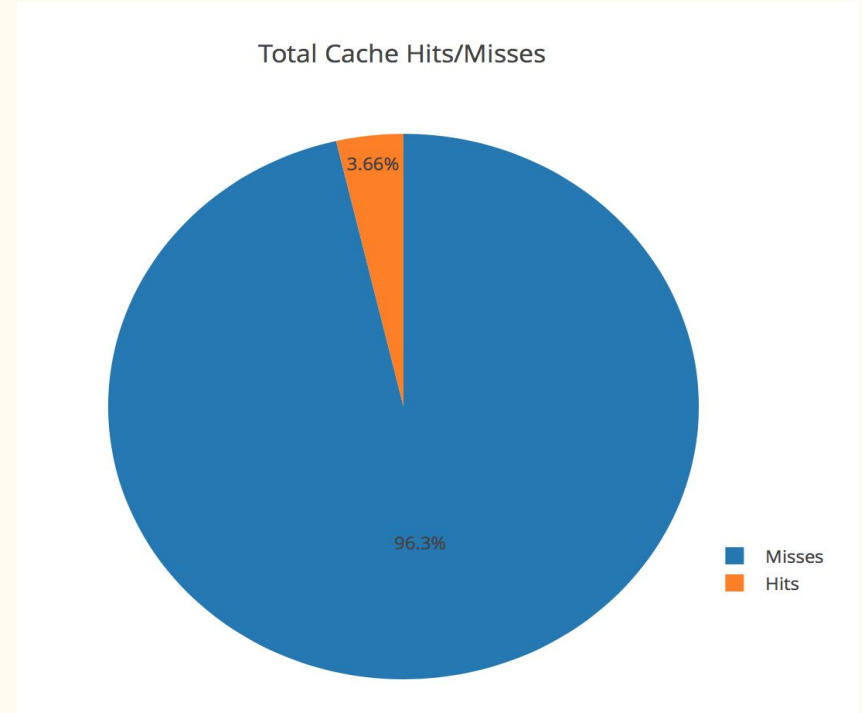
- GHC 46 <ghc46.ghc.andrew.cmu.edu>
- #Physical cores: 8
- L1 data cache: 32KB (#sets = 64, associativity = 8, line size = 64B)

Cache simulator:

- Only simulate L1 data cache
- Write strategy: WRITE_BACK_ALLOCATE
- Coherence protocol: MSI (modified, shared, invalid)
- Interconnect: Directory


Result Analysis

- The overall performance is poor
- Cache hit rate is as low as 3.66%
- BECAUSE each write invalidates all copies of the data



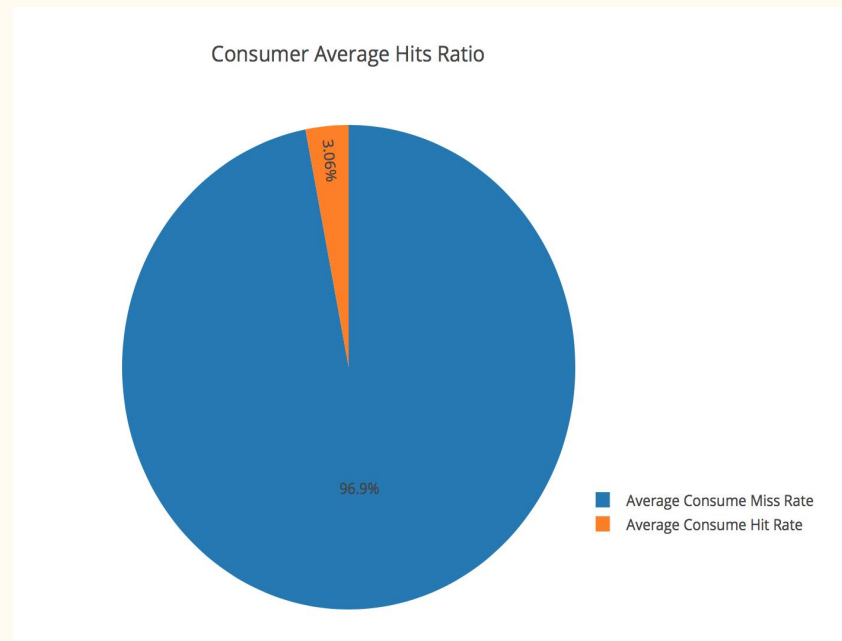
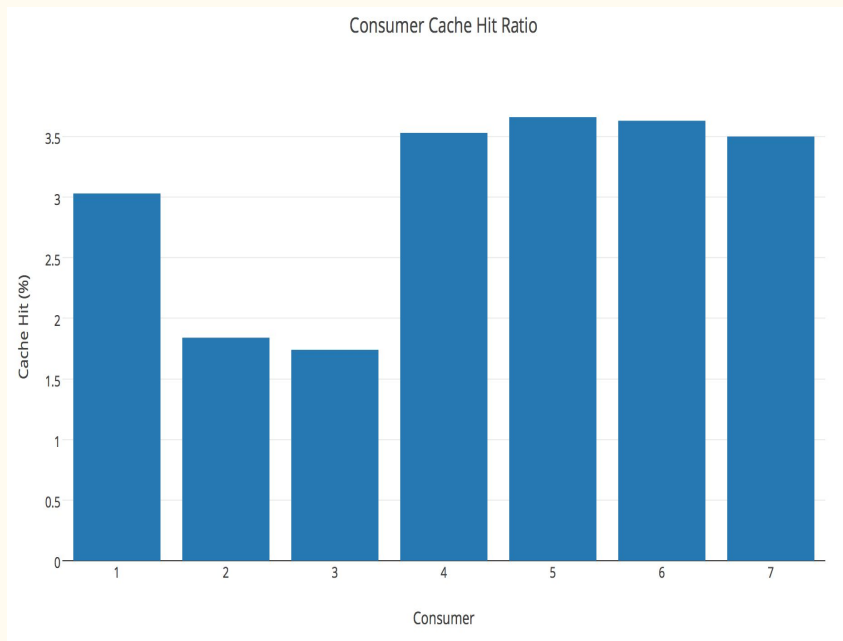
Result Analysis

- Among all hits, the producer thread takes up more than 52%



Result Analysis

Consumer threads have extremely poor cache hit rates



75%



Plans

Plans

TODOs before moving on to 100% goal

1. Refactor cache simulator
2. Support more detailed analysis
 - a. separate load/store hits and misses
 - b. track directory misses
3. Fully test the cache simulator
 - a. deal with unexpected output when the numbers of processors in cache simulator and application are different